

# Chapter 9: Transformation

# Outline

- Transforming the response
  - The Box-Cox method
  - Logit
  - Fisher transformation
- Transforming the predictors
  - Polynomials
  - Regression splines

## Reasons to try transformations

- Nonlinearity
- Non-constant error variance
- Correlated errors
- May improve fit
- Prior information: Incorporate a physical law or some other known relationship

## Examples

$$Y = \exp(\beta_0 + \beta_1 X) \cdot \exp(\epsilon)$$

$$\ln(Y) = \beta_0 + \beta_1 X + \epsilon$$

How to interpret  $\hat{\beta}_1$ ?

$\Rightarrow$  An increase of one in  $X_1$  would multiply the predicted response by  $e^{\hat{\beta}_1}$

## 1.Box-Cox Method

Transformation of the response:  $y \rightarrow g_\lambda(y)$ .

A family of transformations indexed by  $\lambda$  when  $y > 0$ :

$$g_\lambda(y) = \begin{cases} \frac{y^\lambda - 1}{\lambda} & \lambda \neq 0 \\ \ln y & \lambda = 0 \end{cases}$$

## Box-Cox Method Continued

- Can compute **likelihood** of the data using the normal assumption for any given  $\lambda$
- Choose  $\lambda$  to **maximize**:

$$L(\lambda) = -\frac{n}{2} \ln (RSS_{\lambda}/n) + (\lambda - 1) \sum_i \ln y_i$$

- R tries a lot of  $\lambda$ s

## Remarks

RSS: Residual Sum of Square

$$\sum_{i=1}^n (Y_i - \hat{Y}_i)^2$$

A good model has a small RSS

## Box-Cox Method Continued

- In practice, use  $y^\lambda$ .
- In terms of prediction, use maximizer  $\lambda$ .
- In terms of interpretation, use interpretable value near the maximizer  $\lambda$ .



## Box-Cox Method Continued

If  $\hat{\lambda} = 0.46$ , it would be hard to explain what this new response  $y^{0.46}$  means.



If 95% CI for  $\lambda$  contains 0.5,  $\sqrt{y}$  would be preferred because it is easier to interpret.

## Savings & Galapagos Tortoise Examples

Recall from Chapter 4 & 6

```
> library(MASS)
## Box-Cox method for Savings data
> g = lm(sr ~ pop15 + pop75 + dpi + ddpi, savings)
> boxcox(g, plotit=T)
> boxcox(g, plotit=T, lambda=seq(0.5, 1.5, by=0.1))
```

# Savings Example





## Savings & Galapagos Tortoise Examples

Recall from Chapter 4 & 6

```
> library(MASS)
## Box-Cox method for the Tortoise data
> g = lm(Species ~ Area + Elevation + Nearest
+ Scrutz + Adjacent, gala)
> boxcox(g, plotit=T)
> boxcox(g, plotit=T, lambda=seq(0, 1, by=0.05))
```

# Galapagos Tortoise Example



## Transformation in the Tortoise example

```
> summary(lm(Species ~ Area + Elevation + Nearest +  
+           Scruz + Adjacent, data=gala))
```

	Estimate	Std. Error	t value	Pr(> t )
(Intercept)	7.068221	19.154198	0.369	0.715351
Area	-0.023938	0.022422	-1.068	0.296318
Elevation	0.319465	0.053663	5.953	3.82e-06 ***
Nearest	0.009144	1.054136	0.009	0.993151
Scruz	-0.240524	0.215402	-1.117	0.275208
Adjacent	-0.074805	0.017700	-4.226	0.000297 ***

---

Residual standard error: 60.98 on 24 degrees of freedom

Multiple R-squared: 0.7658,      Adjusted R-squared: 0.7171

F-statistic: 15.7 on 5 and 24 DF, p-value: 6.838e-07

## Transformation in the Tortoise example



```
> summary(lm(Species^(1/3) ~ Area + Elevation + Nearest +  
+           Scruz + Adjacent, data=gala))
```

	Estimate	Std. Error	t value	Pr(> t )	
(Intercept)	2.2479224	0.3052013	7.365	1.32e-07	***
Area	-0.0007349	0.0003573	-2.057	0.05070	.
Elevation	0.0054510	0.0008551	6.375	1.37e-06	***
Nearest	0.0118152	0.0167965	0.703	0.48855	
Scruz	-0.0045951	0.0034322	-1.339	0.19317	
Adjacent	-0.0010597	0.0002820	-3.757	0.00097	***

---

Residual standard error: 0.9716 on 24 degrees of freedom  
Multiple R-squared: 0.7543,      Adjusted R-squared: 0.7032  
F-statistic: 14.74 on 5 and 24 DF,   p-value: 1.192e-06

# Diagnostic plots





## Remarks on the Box-Cox Method

- May not choose the  $\lambda$  that exactly maximizes  $L(\lambda)$ , but instead choose one that is **easily interpreted**.
- Sensitive to **outliers**. E.g.,  $\hat{\lambda} = 5$  – ask why?
- If some  $y_i \leq 0$ , can add a constant.
- Transformations of proportions, counts – generalized linear models (later in the course)
- A “quick fix”: if  $y_i$ ’s are **proportions** (range from 0 to 1), consider

$$\ln \left( \frac{y}{1-y} \right)$$

## 2. Logit Method

A special transformations for binomial data (count).

- **Response**  $y_i$ : number of successes out of  $n_i$  independent trials with probability of success  $p_i$

$$\text{Logit}(p) = \log \left( \frac{p}{1-p} \right)$$

## Logit Method: Binomial Data

- $x = (x_1, x_2, \dots, x_p)$ : predictors (quantitative, factors, or both)
- Goal: model the relationship between  $y$  and  $x_1, \dots, x_p$  via modeling the relationship between  $p_i$  and  $x_1, \dots, x_p$ .

## Review: The Binomial Distribution

- $n$  independent trials  $Z_1, \dots, Z_n$
- $P(Z_i = 1) = p$  ("success")  
 $P(Z_i = 0) = 1 - p$  ("failure")
- The binomial variable  $Y = \sum_{i=1}^n Z_i$  is the total number of successes out of  $n$  iid trials


## Review: The Binomial Distribution

- $E(Y) = np$
- $Var(Y) = np(1 - p)$
- Sample proportion (estimate of  $p$ )

$$\hat{p} = \frac{Y}{n}$$

# Logit Method

Logit plot





## Simulation Study

- Y: Response,  $\{0, 1\}$  (e.g. disease ,gender)
- X: Predictor
- Model:

$$\text{Logit}(Y) = \beta_0 + \beta_1 X + \epsilon$$

# Simulation Study





## Simulation Study

```
> glm(y ~ x, family = "binomial")
```

```
Call:  glm(formula = y ~ x, family = "binomial")
```

```
Coefficients:
```

```
(Intercept)                x  
0.4448652914945  0.4014087187817
```

## Transforming the Predictors

Before:

$$y = \beta_0 + \beta_1 x_1 + \cdots + \beta_p x_p + \epsilon$$

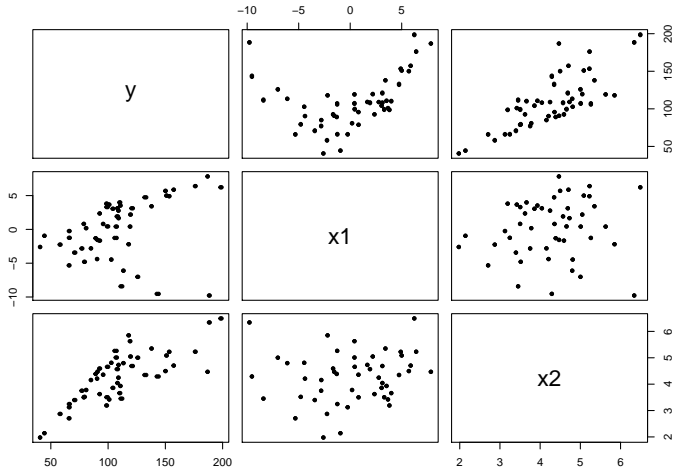
Now:

$$y = \beta_0 + \beta_1 f_1(x) + \cdots + \beta_q f_q(x) + \epsilon$$

$f_j(x)$  are called **basis functions**. Examples:

1. Broken stick regression (skip)
2. Polynomials
3. Regression splines

# Example



## Example

$$Y \propto X_1^2$$

$$Y \propto X_2$$

Transforming predictors is better.

## 2. Polynomials (One Predictor Case)

$$y = \beta_0 + \beta_1 x_1 + \cdots + \beta_d x_1^d + \epsilon$$

How to choose  $d$ :

1. Keep **adding** terms until the new term is not statistically significant
2. Start with a large  $d$  – keep **eliminating** the non-significant highest order term

-Focusing only on **highest** order.

## Polynomials (One Predictor Case) Example




Figure: Scatter Plot

## Forward: Step 1: 1st degree

Coefficients:

Estimate Std. Error t value Pr(>|t|)

(Intercept)	99.614	15.411	6.464	3.99e-09 ***
x	-3.986	1.498	-2.661	0.00911 **

---

Signif. codes: 0 \*\*\* 0.001 \*\* 0.01 \* 0.05 . 0.1 1

Residual standard error: 152.3 on 98 degrees of freedom

Multiple R-squared: 0.06737, Adjusted R-squared: 0.05786

F-statistic: 7.08 on 1 and 98 DF, p-value: 0.009111

## Forward: Step 2: 2nd degree

Coefficients:

	Estimate	Std. Error	t value	Pr(> t )
(Intercept)	-0.0288700	0.1169616	-0.247	0.806
x	0.0104287	0.0098592	1.058	0.293
I(x^2)	1.0006549	0.0006423	1557.883	<2e-16 ***

---  
Signif. codes: 0 \*\*\* 0.001 \*\* 0.01 \* 0.05 . 0.1 1

Residual standard error: 0.9679 on 97 degrees of freedom

Multiple R-squared: 1, Adjusted R-squared: 1

F-statistic: 1.301e+06 on 2 and 97 DF, p-value: < 2.2e-16



## Forward: Step 3: 3rd degree

Coefficients:

	Estimate	Std. Error	t value	Pr(> t )
(Intercept)	-7.916e-02	1.212e-01	-0.653	0.515
x	-9.183e-03	1.652e-02	-0.556	0.580
I(x^2)	1.001e+00	7.363e-04	1359.836	<2e-16 ***
I(x^3)	6.495e-05	4.404e-05	1.475	0.144

---

Signif. codes: 0 '\*\*\*' 0.001 '\*\*' 0.01 '\*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 0.9621 on 96 degrees of freedom

Multiple R-squared: 1, Adjusted R-squared: 1

F-statistic: 8.78e+05 on 3 and 96 DF, p-value: < 2.2e-16

## Backward Elimination

Suppose that the maximum degree  $d = 4$ .

## Backward: Step 1: 4th degree

```
> summary(lm(y ~ x + I(x^2) + I(x^3) + I(x^4) ) )
```

Coefficients:

Estimate Std. Error t value Pr(>|t|)

(Intercept)	2.26e-01	1.55e-01	1.46	0.15
x	-1.79e-02	1.95e-02	-0.92	0.36
I(x^2)	9.97e-01	2.35e-03	423.72	<2e-16 ***
I(x^3)	6.23e-05	7.13e-05	0.87	0.38
I(x^4)	8.22e-06	5.53e-06	1.49	0.14

---

Signif. codes: 0 \*\*\* 0.001 \*\* 0.01 \* 0.05 . 0.1 1

## Backward: Step 2: 3rd degree

Coefficients:

	Estimate	Std. Error	t value	Pr(> t )
(Intercept)	-7.916e-02	1.212e-01	-0.653	0.515
x	-9.183e-03	1.652e-02	-0.556	0.580
I(x^2)	1.001e+00	7.363e-04	1359.836	<2e-16 ***
I(x^3)	6.495e-05	4.404e-05	1.475	0.144

---

Signif. codes: 0 \*\*\* 0.001 \*\* 0.01 \* 0.05 . 0.1 1

Residual standard error: 0.9621 on 96 degrees of freedom

Multiple R-squared: 1, Adjusted R-squared: 1

F-statistic: 8.78e+05 on 3 and 96 DF, p-value: < 2.2e-16

## Backward: Step 3: 2nd degree

Coefficients:

	Estimate	Std. Error	t value	Pr(> t )
(Intercept)	-0.0288700	0.1169616	-0.247	0.806
x	0.0104287	0.0098592	1.058	0.293
I(x^2)	1.0006549	0.0006423	1557.883	<2e-16 ***


---  
Signif. codes: 0 \*\*\* 0.001 \*\* 0.01 \* 0.05 . 0.1 1

Residual standard error: 0.9679 on 97 degrees of freedom

Multiple R-squared: 1, Adjusted R-squared: 1

F-statistic: 1.301e+06 on 2 and 97 DF, p-value: < 2.2e-16

# Issue of Forward Selection



## Issue of Forward Selection

```
> summary(lm(y ~ x ) )
```

Coefficients:

Estimate Std. Error t value Pr(>|t|)

(Intercept)	93.807	13.546	6.93	4.6e-10 ***
-------------	--------	--------	------	-------------

x	-0.303	1.398	-0.22	0.83
---	--------	-------	-------	------

---

Signif. codes: 0 \*\*\* 0.001 \*\* 0.01 \* 0.05 . 0.1 1

## Issue of Forward Selection

```
> summary(lm(y ~ x + I(x^2)) )
```

Coefficients:

Estimate Std. Error t value Pr(>|t|)

(Intercept)	-0.072091	0.119886	-0.60	0.55
x	-0.012966	0.010134	-1.28	0.20
I(x^2)	1.000828	0.000733	1365.37	<2e-16 ***

---

Signif. codes: 0 \*\*\* 0.001 \*\* 0.01 \* 0.05 . 0.1 1

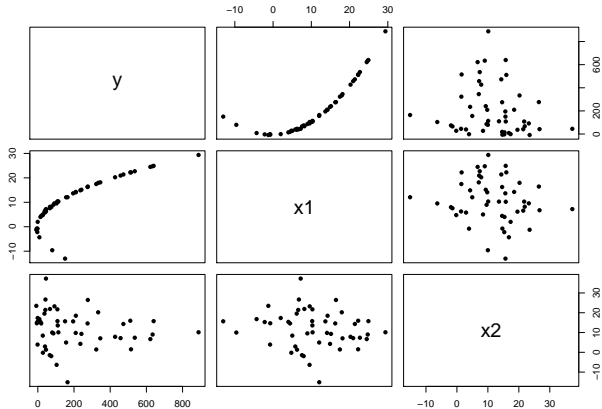


## Polynomials (Two Predictor Case)

- Two Predictors:  $X_1, X_2$
- True Model:

$$Y = 0.5 + X_1 - 0.4X_2 + X_1^2 + \epsilon$$

# Polynomials (Two Predictor Case)



## Forward: Step 1

```
> summary( lm(y~ x1 + x2, data) )$coef
```

```
Estimate Std. Error t value Pr(>|t|)
```

```
(Intercept)    -8.150      34.25  -0.238  8.13e-01
```

```
x1              20.034       1.81  11.095  1.01e-14
```

```
x2             -0.291       1.76  -0.165  8.69e-01
```

## Forward: Step 2

```
> summary( lm(y~ x1 + x2 + I(x1*x2) + I(x1^2) + I(x2^2), data) )
```

	Estimate	Std. Error	t value	Pr(> t )
--	----------	------------	---------	----------

(Intercept)	0.658052	0.217839	3.021	4.19e-03
-------------	----------	----------	-------	----------

x1	0.994548	0.023057	43.135	1.21e-37
----	----------	----------	--------	----------

x2	-0.414351	0.019318	-21.449	6.28e-25
----	-----------	----------	---------	----------

I(x1 * x2)	0.000708	0.001321	0.536	5.95e-01
------------	----------	----------	-------	----------

I(x1^2)	0.999611	0.000607	1645.597	5.19e-107
---------	----------	----------	----------	-----------

I(x2^2)	0.000127	0.000469	0.271	7.88e-01
---------	----------	----------	-------	----------

## Forward: Step 3

```
> summary( lm(y~ x1 + x2 + I(x1*x2) + (x1^2) + I(x2^2)
+ I(x1^2*x2) + I(x1*x2^2) + I(x1^3) + I(x2^3), data) )$coef
```

	Estimate	Std. Error	t value	Pr(> t )
(Intercept)	-67.62499	15.40574	-4.39	7.79e-05
x1	18.97354	1.63939	11.57	1.70e-14
x2	13.40205	2.37122	5.65	1.35e-06
I(x1 * x2)	-1.84729	0.21621	-8.54	1.20e-10
I(x2^2)	-0.59440	0.12344	-4.82	2.03e-05
I(x1^2 * x2)	0.04707	0.00409	11.50	2.07e-14
I(x1 * x2^2)	0.03513	0.00871	4.03	2.34e-04
I(x1^3)	0.01457	0.00202	7.22	8.09e-09
I(x2^3)	0.00671	0.00159	4.22	1.34e-04

## Forward: Step 4

```
> summary( lm(y~ x1 + x2 + I(x1*x2) + (x1^2) + I(x2^2) + I(x1^2*x2) + I(x1*x2^2) + I(x1^3) + I(x2^3)
+           + I(x1^3*x2) + I(x1*x2^3)+ I(x1^2*x2^2) + I(x1^4) + I(x2^4)
+           , data) )$coef
```

Estimate Std. Error t value Pr(>|t|)

(Intercept)	-3.60e+01	6.72e+00	-5.36	5.04e-06
x1	1.23e+01	9.86e-01	12.46	1.29e-14
x2	1.25e+01	1.56e+00	7.98	1.78e-09
I(x1 * x2)	-2.24e+00	1.89e-01	-11.84	5.62e-14
I(x2^2)	-1.05e+00	1.40e-01	-7.53	6.65e-09
I(x1^2 * x2)	1.03e-01	4.27e-03	24.14	7.99e-24
I(x1 * x2^2)	1.17e-01	1.26e-02	9.28	4.43e-11
I(x1^3)	3.46e-02	2.38e-03	14.50	1.32e-16
I(x2^3)	2.90e-02	4.73e-03	6.12	4.80e-07
I(x1^3 * x2)	-1.56e-03	1.26e-04	-12.39	1.53e-14
I(x1 * x2^3)	-1.71e-03	2.83e-04	-6.06	5.87e-07
I(x1^2 * x2^2)	-2.30e-03	2.65e-04	-8.70	2.24e-10
I(x1^4)	-3.86e-04	6.04e-05	-6.40	2.05e-07
I(x2^4)	-2.37e-04	4.73e-05	-5.02	1.42e-05

## Issue

Finding an optimal  $d$  may be impossible

## Savings Example: Forward Selection

```
# tired of typing data = savings?  
> attach(savings)
```

```
## Polynomials
```

```
## 1st degree
```

```
> summary(lm(sr ~ ddpi))
```

Coefficients:

	Estimate	Std.Error	t value	Pr(> t )
(Intercept)	7.8830	1.0110	7.797	4.46e-10
ddpi	0.4758	0.2146	2.217	0.0314



```
## 2nd degree
```

```
> summary(lm(sr ~ ddpi + I(ddpi^2)))
```

	Estimate	Std.Error	t value	Pr(> t )
(Intercept)	5.13038	1.43472	3.576	0.000821
ddpi	1.75752	0.53772	3.268	0.002026
I(ddpi^2)	-0.09299	0.03612	-2.574	0.013262

```
## 3rd degree
```

```
> summary(lm(sr ~ ddpi + I(ddpi^2) + I(ddpi^3)))
```

	Estimate	Std.Error	t value	Pr(> t )
Intercept	5.145e+00	2.199e+00	2.340	0.0237
ddpi	1.746e+00	1.380e+00	1.265	0.2123
ddpi^2	-9.097e-02	2.256e-01	-0.403	0.6886
ddpi^3	-8.497e-05	9.374e-03	-0.009	0.9928

## Linear Transformation

Linear Transformation of a predictor does **not** change its p-value.

```
## Be careful with elimination
```

```
> mddpi = ddpi - 10
```

```
> summary(lm(sr ~ mddpi + I(mddpi^2)))
```

Coefficients:

	Estimate	Std.Error	t value	Pr(> t )
Intercept	13.40705	1.42401	9.415	2.16e-12
mddpi	-0.10219	0.30274	-0.338	0.7372
mddpi^2	-0.09299	0.03612	-2.574	0.0133

## Orthogonal Polynomials

For numerical stability:

$$z_1 = a_1 + b_1x$$

$$z_2 = a_2 + b_2x + c_2x^2$$

$$z_3 = a_3 + b_3x + c_3x^2 + d_3x^3$$

$$\vdots = \vdots$$

$a, b, c \dots$  are chosen so that  $z_j^T z_{j'} = 0$  when  $j \neq j'$ .

## Savings Example

```
## Orthogonal polynomials
```

```
> summary(lm(sr ~ poly(ddpi, 4)))
```

Coefficients:

	Estimate	Std. Error	t value	Pr(> t )
(Intercept)	9.67100	0.58460	16.543	<2e-16 ***
poly(ddpi, 4)1	9.55899	4.13376	2.312	0.0254 *
poly(ddpi, 4)2	-10.49988	4.13376	-2.540	0.0146 *
poly(ddpi, 4)3	-0.03737	4.13376	-0.009	0.9928
poly(ddpi, 4)4	3.61197	4.13376	0.874	0.3869

Residual standard error: 4.134 on 45 degrees of freedom

Multiple R-Squared: 0.2182      Adjusted R-squared: 0.1488

F-statistic: 3.141 on 4 and 45 DF      p-value: 0.02321

## Polynomials in several predictors

Define polynomials in more than one variable. E.g.

$$y = \beta_0 + \beta_1 x_1 + \beta_2 x_2 + \beta_{11} x_1^2 + \beta_{22} x_2^2 + \beta_{12} x_1 x_2 + \epsilon$$


R command:

```
> g = lm(sr ~ poly(pop15, ddpi, degree=2))
```

# Broken Stick Regression

Sometimes, there are different linear regression models apply in different regions of the data.

- Economic Crisis



## Broken Stick Regression: Basis

An important property of broken stick regression is **continuity**.

Hence we define the following basis functions:

$$B_l(x) = \begin{cases} c - x & \text{if } x < c \\ 0 & \text{otherwise} \end{cases} \quad B_r(x) = \begin{cases} x - c & \text{if } x \geq c \\ 0 & \text{otherwise} \end{cases}$$

where  $c$  marks the division between the two groups.

$B_l$  and  $B_r$  form a **first-order spline basis** with a knotpoint at  $c$ .



## Broken Stick Regression Model

The Form of the model:

$$y = \beta_0 + \beta_1 B_l + \beta_2 B_r + \epsilon.$$

Example:

```
> lhs <- function (x) ifelse(x < 35, 35!x, 0)
> rhs <- function (x) ifelse(x < 35, 0, x!35)
> gb <- lm (sr ~ lhs (pop15) + rhs(pop15), savings)
> x <- seq(20, 48, by=1)
> py <- gb$coef[1]+gb$coef[2]*lhs(x)+gb$coef[3]*rhs(x)
> lines (x, py, lty=2)
```

### 3. Regression Splines

Disadvantage of polynomials: each data point affects the fit globally. Remedy: *B-spline*.

*Cubic B-spline* basis functions on interval  $(a, b)$  with pre-specified knots  $t_1, \dots, t_k$ :

- Non-zero on interval defined by four successive knots and zero elsewhere  $\Rightarrow$  *local* influence property
- Cubic polynomial fit to each four successive knots
- *Smooth*
- Integrates to one

## Simulation Example

$$y = \sin^3(2\pi x^3) + \epsilon, \quad \epsilon \sim N(0, 0.1^2)$$

- Not a polynomial, not a cubic spline...
- But smooth and has many inflection points


## Simulation Example

```
> ## Data generation
> myf = function(x) sin(2*pi*x^3)^3
> x = seq(0, 1, by=0.01)
> y = myf(x) + 0.1*rnorm(101)
> matplot(x, cbind(y, myf(x)), type="pl", lwd= 3, pch = 20,
  main ="True Function")


> ## Polynomials
> g4 = lm(y ~ poly(x, 4))
> g12 = lm(y ~ poly(x, 12))
> matplot(x, cbind(y, g4$fit, g12$fit), type="plll", pch = 20, lty=1,
  lwd= 3, col = c("black","red","green"), main="Polynomial" )
```

# Polynomial results

True Function




Polynomial





```
## Regression splines
> library(splines)
> knots = c(0, 0, 0, 0, 0.2, 0.4, 0.5, 0.6,
            0.7, 0.8, 0.85, 0.9, 1, 1, 1, 1)
> bx = splineDesign(knots, x)
> gs = lm(y ~ bx)
> matplot(x, bx, type="l")
> matplot(x, cbind(y, gs$fit), type="pl")
```

# Spline results



## 4. Smoothing Splines





## Other Transformations

- Smoothing splines
- Generalized additive models
- CART, MARS, MART, **neural networks**

Rule of thumb:

- for large data sets, complex models are better (**with appropriate control of the number of parameters**);
- for small data sets or high noise levels (e.g., social sciences), standard regression is more appropriate.

## Other Transformations

As we have a better computer,

New Rule of thumb:

Mixtures of complex models are better **with appropriate control of the number of parameters**;