

Chapter 10: Shrinkage Methods

Outline

- Principal Components Analysis
- Ridge regression
- Lasso
- Partial Least Squares

Training and Test Data Set


Motivation

- Model Selections and Shrinkage Methods provide good models
- Hard to choose only one optimal model.
- Choose an optimal model based on its performance.

Training and Test Data Set

Method


- Split dataset into two groups .
- **Training Set**: Used to train the model.
- **Test Set**: Used to measure the performance of the trained model.
- **Validation Set**: Used to train tuning parameters. (Extra).



Training and Test Data Set

Why **Test Set**?

- **Overfitting** Issue: It refers to a model that models the training data too well.



Training and Test Data Set

Why Test Set?

- Overfitting Issue
- Complex models tend to have a good performance in training data.
- A good model in training dataset may not be a good model in new dataset.

Training and Test Data Set

How to determine Test Data Set?

- Absolute Random selection: Choose 10 ~ 20% of data set.
- Choose 10 ~ 20% of data set with same proportion of success in both Training and Test data set. (To prevent the test set only contains success or same value of predictors.)

Training and Test Data Set

Weakness

- Cannot be performed when sample size is **small**
- Cross-Validation (Later)


Principal Components Analysis (PCA)

- Special transformation on predictors
- Useful for high dimensional data
- Solve collinearity issue

Principal Components Analysis (PCA)

- Find the u_1 such that $\text{var}(u_1^T X)$ is maximized subject to $u_1^T u_1 = 1$.
- Find the u_2 such that $\text{var}(u_2^T X)$ is maximized subject to $u_2^T u_1 = 0$ and $u_2^T u_2 = 1$.
- Keep finding directions of greatest variation orthogonal to those directions we have already found.

Principal Components Analysis (PCA)



Principal Components Analysis (PCA)




Figure: PCA: u_1 and u_2

Simulation Study

```
> prcomp(x)
```

Rotation:



	PC1	PC2
x1	-0.892	-0.451
x2	-0.451	0.892

$$Z_1 = -0.892X_1 - 0.451X_2$$

$$Z_2 = -0.451X_1 + 0.892X_2$$

Simulation Study

```
> plot(x, pch = 16, xlim = c(-4, 4), ylim = c(-3,3))  
> abline(a = 0, b = 0.451/-0.892, col = "red")  
> abline(a = 0, b = 0.451/0.892, col = "blue")
```



Simulation Study

```
> prX = prcomp(x)
> summary(prX)
```

Importance of components:

	PC1	PC2
Standard deviation	1.210	0.1918
Proportion of Variance	0.976	0.0245
Cumulative Proportion	0.976	1.0000

Simulation Study

```
> round(prX$rot[,1],3)
```

```
x1      x2
```

```
-0.892 -0.451
```

$$Z_1 = -0.892X_1 - 0.451X_2$$

Z_1 explains 97.6% of the both X_1 and X_2

Simulation Study

```
> lm0 = lm(Y ~ X1 + X2)
> summary(lm0)
```

Coefficients:

Estimate Std. Error t value Pr(>|t|)

(Intercept)	-0.1164	0.1091	-1.067	0.289	
x1	0.8597	0.2020	4.255	4.82e-05	***
x2	1.2688	0.2219	5.717	1.20e-07	***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 1.086 on 97 degrees of freedom

Multiple R-squared: 0.7988, Adjusted R-squared: 0.7946

F-statistic: 192.5 on 2 and 97 DF, p-value: < 2.2e-16

Simulation Study

```
> lm1 = lm(Y ~ Z)
> summary(lm1)
```

Coefficients:

Estimate Std. Error t value Pr(>|t|)

(Intercept)	-0.1092	0.1094	-0.999	0.32
z	-1.4875	0.0762	-19.520	<2e-16 ***

Signif. codes: 0 *** 0.001 ** 0.01 * 0.05 . 0.1 1

Residual standard error: 1.09 on 98 degrees of freedom

Multiple R-squared: 0.7954, Adjusted R-squared: 0.7933

F-statistic: 381 on 1 and 98 DF, p-value: < 2.2e-16

Fat Example

- Response: fat
- Predictors: neck, chest, abdom, hip, thigh, knee, ankle, biceps, forearm, wrist

```
> cfat = fat[,9:18]
> prfat = prcomp(cfat)
> dim(prfat$rot)
[1] 10 10
> dim(prfat$x)
[1] 252 10
```

Fat Example

```
> summary(prfat)
```

Importance of components:

	PC1	PC2	PC3	PC4	PC5	PC6	PC7	PC8	PC9	PC10
Standard deviation	15.990	4.0658	2.9660	2.0004	1.69408	1.49881	1.30322	1.25478	1.10955	0.52737
Proportion of Variance	0.867	0.0561	0.0298	0.0136	0.00973	0.00762	0.00576	0.00534	0.00417	0.00094
Cumulative Proportion	0.867	0.9230	0.9529	0.9664	0.97617	0.98378	0.98954	0.99488	0.99906	1.00000

Fat Example

```
> round(prfat$rot[,1],2)
```

neck	chest	abdom	hip	thigh	knee
0.12	0.50	0.66	0.42	0.28	0.12

ankle	biceps	forearm	wrist
0.06	0.15	0.07	0.04

Fat Example

```
> prfatc = prcomp(cfat, scale = TRUE)
> summary(prfatc)
```

Importance of components:

	PC1	PC2	PC3	PC4	PC5	PC6	PC7	PC8	PC9	PC10
Standard deviation	2.650	0.8530	0.8191	0.7011	0.5471	0.5283	0.4520	0.4054	0.27827	0.2530
Proportion of Variance	0.702	0.0728	0.0671	0.0492	0.0299	0.0279	0.0204	0.0164	0.00774	0.0064
Cumulative Proportion	0.702	0.7749	0.8420	0.8911	0.9211	0.9490	0.9694	0.9859	0.99360	1.0000

Fat Example

```
> round(prfatc$rot[,1],3)
```

neck	chest	abdom	hip	thigh	knee	ankle
0.327	0.339	0.334	0.348	0.333	0.329	0.247

biceps	forearm	wrist
0.322	0.270	0.299

Fat Example

```
> round(prfatc$rot[,2],3)
```

neck	chest	abdom	hip	thigh	knee	ankle
-0.003	-0.273	-0.398	-0.255	-0.191	0.022	0.625
biceps	forearm	wrist				
0.022	0.363	0.377				

Fat Example

```
> lmoda = lm(fat$brozek ~., data = cfat)
```

```
> summary(lmoda)
```

Coefficients:

Estimate Std. Error t value Pr(>|t|)

(Intercept)	7.228749	6.214309	1.163	0.24588	
neck	-0.581947	0.208580	-2.790	0.00569	**
chest	-0.090847	0.085430	-1.063	0.28866	
abdom	0.960229	0.071582	13.414	< 2e-16	***
hip	-0.391355	0.112686	-3.473	0.00061	***
thigh	0.133708	0.124922	1.070	0.28554	
knee	-0.094055	0.212394	-0.443	0.65828	
ankle	0.004222	0.203175	0.021	0.98344	
biceps	0.111196	0.159118	0.699	0.48533	
forearm	0.344536	0.185511	1.857	0.06450	.
wrist	-1.353472	0.471410	-2.871	0.00445	**

Residual standard error: 4.071 on 241 degrees of freedom

Multiple R-squared: 0.7351, Adjusted R-squared: 0.7241

F-statistic: 66.87 on 10 and 241 DF, p-value: < 2.2e-16

Fat Example

```
> lmodpcr = lm(fat$brozek ~ prfatc$x[,1:2])
```

```
> summary(lmodpcr)
```

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
--	----------	------------	---------	----------

(Intercept)	18.9385	0.3291	57.542	<2e-16 ***
-------------	---------	--------	--------	------------

prfatc\$x[, 1:2]PC1	1.8420	0.1245	14.800	<2e-16 ***
---------------------	--------	--------	--------	------------

prfatc\$x[, 1:2]PC2	-3.5505	0.3866	-9.184	<2e-16 ***
---------------------	---------	--------	--------	------------

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 5.225 on 249 degrees of freedom

Multiple R-squared: 0.5492, Adjusted R-squared: 0.5456

F-statistic: 151.7 on 2 and 249 DF, p-value: < 2.2e-16

Fat Example

```
> lmodpcr2 = lm(fat$brozek ~ prfatc$x[,1:10])  
> summary(lmodpcr2)
```

Estimate	Std. Error	t value	Pr(> t)
(Intercept)	18.93849	0.25647	73.843 < 2e-16 ***
prfatc\$x[, 1:10]PC1	1.84198	0.09698	18.993 < 2e-16 ***
prfatc\$x[, 1:10]PC2	-3.55053	0.30126	-11.785 < 2e-16 ***
prfatc\$x[, 1:10]PC3	0.25669	0.31374	0.818 0.414067
prfatc\$x[, 1:10]PC4	0.54094	0.36652	1.476 0.141273
prfatc\$x[, 1:10]PC5	3.72632	0.46973	7.933 8.03e-14 ***
prfatc\$x[, 1:10]PC6	-1.48784	0.48642	-3.059 0.002474 **
prfatc\$x[, 1:10]PC7	1.94878	0.56859	3.427 0.000716 ***
prfatc\$x[, 1:10]PC8	-0.12247	0.63390	-0.193 0.846967
prfatc\$x[, 1:10]PC9	-1.71366	0.92351	-1.856 0.064731 .
prfatc\$x[, 1:10]PC10	-9.01059	1.01566	-8.872 < 2e-16 ***

Residual standard error: 4.071 on 241 degrees of freedom

Multiple R-squared: 0.7351, Adjusted R-squared: 0.7241

F-statistic: 66.87 on 10 and 241 DF, p-value: < 2.2e-16

Fat Example

Benefits of PCA

- Orthogonal Predictors
- No collinearity Issue
- Sometimes easy to Interpret

Choice of the number of variables

How to choose the optimal number of variables.

- **Interpretability**: It is important to examine the interpretability of the components and make sure that those providing a interpretable result are retained.
- **Total variance**
- Eigenvalues (Skip)

Remarks on PCA

- Interpretation may be easy or difficult
- Sufficiently reduce the number of predictors
- Difficult to decide the number of predictors

Choice of the number of variables

How to choose the optimal number of variables.

- Training and Test Set
- Root mean square of error (RMSE)

Root mean square of error (RMSE)


Definition:

$$\text{RMSE} = \sqrt{\frac{1}{n} \sum_i^n (\hat{y}_i - y_i)^2}$$

Motivation:

- Bias: $y - \mathbb{E}(\hat{y})$
- Variance: $\text{Var}(\hat{y})$

Bias



Variance

$$y_1 = 10$$

- Small variance case: 95% confidence interval for y_1 :
(9.99, 10.01)
→ 95% sure that y_1 is in (9.99, 10.01)
- :Large variance case: 95% confidence interval for y_1 : (-5, 25)
→ 95% sure that y_1 is in (-5, 25)

Root mean square of error (RMSE)

- Mean Square Error = $\text{Bias}(\hat{y})^2 + \text{Variance}(\hat{y})$
- RMSE is a good criterion to choose an optimal model

Simulation Study: Choice of the number of variables

How to choose the optimal number of variables.

- 4 predictors: X_1, X_2, X_3, X_4
- Training set: 40 observations
- Test set: 10 observations

Simulation Study

```
> ran = sample(1:50, replace = F)[1:40]
> train = data[ran,]
> test = data[setdiff(1:50,ran),]
> prx = prcomp(x[ran,])
> summary(prx)
```

Importance of components:

	PC1	PC2	PC3	PC4
Standard deviation	1.3194	0.5817	0.4301	0.21264
Proportion of Variance	0.7538	0.1465	0.0801	0.01958
Cumulative Proportion	0.7538	0.9003	0.9804	1.00000

Simulation Study

```
> lmodpca2 = lm(y ~ prx$x[,1:2], train)
> z1 = prx$rotation[,1] %*% t( test[,1:4] )
> z2 = prx$rotation[,2] %*% t( test[,1:4] )
> z = rbind(z1, z2)
> ypred = coef(lmodpca2) %*% rbind(1, z)
> sqrt( mean( (test$y - ypred)^2) )
[1] 1.294111
```

Simulation Study

```
> lmodpca3 = lm(y ~ prx$x[,1:3], train)
> z1 = prx$rotation[,1] %*% t( test[,1:4] )
> z2 = prx$rotation[,2] %*% t( test[,1:4] )
> z3 = prx$rotation[,3] %*% t( test[,1:4] )
> z = rbind(z1, z2, z3)
> ypred = coef(lmodpca3) %*% rbind(1, z)
> sqrt( mean( (test$y - ypred)^2) )
[1] 1.29453
```

Ridge Regression

Penalizing the square of the coefficients

$$\min_{\beta} \sum_{i=1}^n \left(y_i - \beta_0 - \sum_{j=1}^p \beta_j x_{ij} \right)^2 + \lambda \sum_{j=1}^p \beta_j^2$$

Assumption:

- Regression Coefficients should not be very large (after standardization).
- A large number of predictors should be considered.
- High collinearity exists.

Ridge Regression

Penalizing the square of the coefficients

$$\min_{\beta} \sum_{i=1}^n \left(y_i - \beta_0 - \sum_{j=1}^p \beta_j x_{ij} \right)^2 + \lambda \sum_{j=1}^p \beta_j^2$$


- The coefficients $\hat{\beta}^{\text{ridge}}$ are shrunk towards zero.
- $\lambda \geq 0$ is a **tuning parameter**.
- λ controls the amount of shrinkage.
- What happens if $\lambda \rightarrow 0$?
- What happens if $\lambda \rightarrow \infty$?

Equivalent Formulation

$$\begin{aligned} \min_{\beta} \quad & \sum_{i=1}^n \left(y_i - \beta_0 - \sum_{j=1}^p \beta_j x_{ij} \right)^2 \\ \text{subject to} \quad & \sum_{j=1}^p \beta_j^2 \leq s \end{aligned}$$

- Explicitly constraint the size of the coefficients.

Equivalent Formulation



Ridge Regression

When there are many **highly correlated variables**

- $\hat{\beta}^{\text{ols}}$ may have a large coefficient on one variable and a similarly large negative coefficient on its correlated variable (**Unstable**).
- In ridge regression, the size constraint tries to avoid this phenomenon.

Often standardize the predictors first.

Solution for Ridge Regression

- The solution is

$$\hat{\beta}^{\text{ridge}} = (\mathbf{X}^T \mathbf{X} + \lambda \mathbf{I})^{-1} \mathbf{X}^T \mathbf{y}$$

- $\hat{\beta}$ is linear in \mathbf{y} .
- $\hat{\beta}$ is biased.

Comparison to LSE

$$\hat{\beta}^{\text{ridge}} = (\mathbf{X}^T \mathbf{X} + \lambda \mathbf{I})^{-1} \mathbf{X}^T \mathbf{y}$$

- Even if \mathbf{X} is not full-rank, $(\mathbf{X}^T \mathbf{X} + \lambda \mathbf{I})$ is invertible, thus solve exact collinearity issue.
- $\hat{\beta}^{\text{ridge}}$ has smaller variance than the OLS, thus may have smaller mean square error (MSE).

Shrinkage in Ridge

Suppose **orthonormal design** ($\mathbf{X}^\top \mathbf{X} = \mathbf{I}$). Then $\hat{\beta}^{\text{ols}} = \mathbf{X}^\top \mathbf{y}$, and

$$(\mathbf{y} - \mathbf{X}\beta)^\top (\mathbf{y} - \mathbf{X}\beta) = \text{constant} + \sum_{j=1}^p (\beta_j - \hat{\beta}_j^{\text{ols}})^2.$$

Then ridge regression minimizes

$$\sum_{j=1}^p (\beta_j - \hat{\beta}_j^{\text{ols}})^2 + \lambda \sum_{j=1}^p \beta_j^2.$$

Equivalent to the component-wise minimization

$$\min_{\beta_j} (\beta_j - \hat{\beta}_j^{\text{ols}})^2 + \lambda \beta_j^2 \implies \hat{\beta}_j^{\text{ridge}} = \frac{1}{1 + \lambda} \hat{\beta}_j^{\text{ols}}.$$

Shrinkage in Ridge

- Shrink the estimate towards zero by a positive constant less than 1
- $\text{Var}(\hat{\beta}_j^{\text{ridge}}) = \frac{1}{(1+\lambda)^2} \text{Var}(\hat{\beta}_j^{\text{ols}})$.
- $\lambda \uparrow$, shrinkage \uparrow , bias \uparrow , variance \downarrow
- $\lambda \downarrow$, shrinkage \downarrow , bias \downarrow , variance \uparrow .

Simulation Study: Almost Independent Predictors

3 predictors: X_1, X_2, X_3

```
> cor(data[,1:3])
```

	x1	x2	x3
x1	1.0000	0.0505	-0.142
x2	0.0505	1.0000	-0.143
x3	-0.1425	-0.1428	1.000

```
> lmod = lm(y ~ x1 + x2 + x3, data)
```

```
> lmod$coefficients
```

(Intercept)	x1	x2	x3
-0.724	1.023	1.299	1.817

Simulation Study

```
> require(MASS)
> #lambda = 0
> lmrid = lm.ridge(y~x1 + x2 + x3, data, lmabda = 0)
> lmrid
```

(Intercept)	x1	x2	x3
-0.724	1.023	1.299	1.817

Simulation Study

```
> #lambda = 0.1
> lmrid2 = lm.ridge(y~x1 + x2 + x3, data, lambda = 0.1)
> lmrid2
```

(Intercept)	x1	x2	x3
-0.669	1.021	1.297	1.809

```
>
> #lambda = 10
> lmrid3 = lm.ridge(y~x1 + x2 + x3, data, lambda = 10)
> lmrid3
```

(Intercept)	x1	x2	x3
3.819	0.848	1.116	1.187

Simulation Study: Correlated Predictors

```
> lmod$coefficients
```

(Intercept)	x1	x2	x3
-0.676	1.132	-0.275	1.779

```
> cor(data[,1:3])
```

	x1	x2	x3
x1	1.000	0.996	0.999
x2	0.996	1.000	0.998
x3	0.999	0.998	1.000

Simulation Study: Correlated Predictors

```
> require(MASS)
> #lambda = 0
> lmrid = lm.ridge(y~x1 + x2 + x3, data, lmabda = 0)
> lmrid
```

(Intercept)	x1	x2	x3
-0.676	1.132	-0.275	1.779

```
>
> #lambda = 0.1
> lmrid2 = lm.ridge(y~x1 + x2 + x3, data, lambda = 0.1)
> lmrid2
```

(Intercept)	x1	x2	x3
-0.720	1.405	0.327	1.195

Simulation Study: Correlated Predictors

```
> #lambda = 1
> lmrid3 = lm.ridge(y~x1 + x2 + x3, data, lambda = 1)
> lmrid3
```

(Intercept)	x1	x2	x3
-0.166	1.260	0.968	0.848

```
>
> #lambda = 10
> lmrid3 = lm.ridge(y~x1 + x2 + x3, data, lambda = 10)
> lmrid3
```

(Intercept)	x1	x2	x3
4.85	1.12	1.07	0.74

Simulation Study: Comparison to LSE

```
# Generate Training/Test sets
> ran = sample(1:50, replace = F)[1:40]
> train = data[ran,]
> test = data[setdiff(1:50,ran),]

> lmod = lm(y ~ x1 + x2 + x3, train)
> lmrid = lm.ridge(y~x1 + x2 + x3, train, lambda = 0.1)
> lmrid2 = lm.ridge(y~x1 + x2 + x3, train, lambda = 1)
```

Simulation Study: Comparison to LSE

```
# RMSE
```

```
> sqrt(mean((test$y - predict(lmod,test))^2))
```

```
[1] 5.4921
```

```
>
```

```
> ypred = cbind(1, as.matrix(test[,-4])) %*% coef(lmrid)
```

```
> sqrt(mean((test$y - ypred)^2))
```

```
[1] 5.1544
```

```
>
```

```
> ypred = cbind(1, as.matrix(test[,-4])) %*% coef(lmrid2)
```

```
> sqrt(mean((test$y - ypred)^2))
```

```
[1] 5.5591
```


Simulation Study: Normalization

```
> data.scale = scale(data[,1:3])
> data.scale = data.frame(data.scale, y = data$y)
>
> train = data.scale[ran,]
> test = data.scale[setdiff(1:50,ran),]
>
> lmod = lm(y ~ x1 + x2 + x3, train)
> lmrid = lm.ridge(y~x1 + x2 + x3, train, lambda = 0.1)
> lmrid2 = lm.ridge(y~x1 + x2 + x3, train, lambda = 1)
```

Simulation Study: Normalization

```
> sqrt(mean((test$y - predict(lmod,test))^2))  
[1] 4.8518  
>  
> ypred = cbind(1, as.matrix(test[, -4])) %*% coef(lmrid)  
> sqrt(mean((test$y - ypred)^2))  
[1] 4.8484  
>  
> ypred = cbind(1, as.matrix(test[, -4])) %*% coef(lmrid2)  
> sqrt(mean((test$y - ypred)^2))  
[1] 4.8587
```

Choice of Tuning Parameter λ

Determination of the tuning parameter λ

- **Generalized Cross-Validation** (GCV): Almost same concept as Training and Test Set.

- GCV

$$V(\lambda) = \frac{\frac{1}{n} \|(I - A(\lambda))y\|^2}{\left(\frac{1}{n} \text{tr}(I - A(\lambda))\right)^2},$$

where $A(\lambda) = X(X^T X + n\lambda I)^{-1} X^T$.

- R provides a good tuning parameter λ (Not always optimal).

Simulation Study

```
> lmrid = lm.ridge(y~x1 + x2 + x3, train, lambda = seq(0,1, len
> lmrid$GCV
0.00      0.25      0.50      0.75      1.00
0.68954 0.66428 0.66800 0.67131 0.67449
>
> which.min(lmrid$GCV)
0.25
```

Simulation Study

```
> lmrid = lm.ridge(y~x1 + x2 + x3, train,  
                  lambda = seq(0,1, len = 100))  
> which.min(lmrid$GCV)  
0.161616  
  
> lmrid_GCV = lm.ridge(y~x1 + x2 + x3, train, lambda = 0.161616)  
> ypred = cbind(1, as.matrix(test[,-4])) %*% coef(lmrid_GCV)  
> sqrt(mean((test$y - ypred)^2))  
[1] 4.8405
```

Simulation Study: Another Training/Test Set

```
> ran = sample(1:50, replace = F)[1:40]
> train = data.scale[ran,]
> test = data.scale[setdiff(1:50,ran),]
>
> lmrid = lm.ridge(y~x1 + x2 + x3, train,
                  lambda = seq(0,1, len = 100))
> which.min(lmrid$GCV)
0.080808
```

Simulation Study: Another Training/Test Set

```
> ypred = cbind(1, as.matrix(test[,-4])) %*% coef(lmrid_GCV)
> lmrid_GCV = lm.ridge(y~x1 + x2 + x3, train, lambda = 0.0808)
>
> ypred = cbind(1, as.matrix(test[,-4])) %*% coef(lmrid_GCV)
> sqrt(mean((test$y - ypred)^2))
[1] 4.486
>
> lmod = lm(y ~ x1 + x2 + x3, train)
> sqrt(mean((test$y - predict(lmod,test))^2))
[1] 4.7677
```

Simulation Study: Comparison to PCA

```
> X = train[,1:3]
```

```
> prx = prcomp(X)
```

```
> summary(prx)
```

Importance of components:

	PC1	PC2	PC3
--	-----	-----	-----

Standard deviation	1.755	0.06021	0.03225
--------------------	-------	---------	---------

Proportion of Variance	0.998	0.00117	0.00034
------------------------	-------	---------	---------

Cumulative Proportion	0.998	0.99966	1.00000
-----------------------	-------	---------	---------

Simulation Study: Comparison to PCA

```
> z = 0.57778 * X[,1] + 0.57848 * X[,2] + 0.57579 * X[,3]
> lmodpcr = lm(train$y ~ z)
> summary(lmodpcr)
```

Coefficients:

Estimate Std. Error t value Pr(>|t|)

(Intercept)	88.530	0.852	103.9	<2e-16 ***
-------------	--------	-------	-------	------------

z	29.304	0.491	59.6	<2e-16 ***
---	--------	-------	------	------------

Signif. codes: 0 *** 0.001 ** 0.01 * 0.05 . 0.1 1

Residual standard error: 5.39 on 38 degrees of freedom

Multiple R-squared: 0.989, Adjusted R-squared: 0.989

F-statistic: 3.56e+03 on 1 and 38 DF, p-value: <2e-16

Simulation Study: Comparison to PCA

```
> ypred = cbind(1, as.matrix(0.57778 * test[,1]
+ 0.57848 * test[,2] + 0.57579 * test[,3]))
%% coef(lmodpcr)
> sqrt(mean((test$y - ypred)^2))
[1] 4.1143
```

LASSO


Least absolute shrinkage and selection operator (Chen, Donoho and Saunders 1996; Tibshirani 1996)

$$\min_{\beta} \sum_{i=1}^n \left(y_i - \beta_0 - \sum_{j=1}^p \beta_j x_{ij} \right)^2 + \lambda \sum_{j=1}^p |\beta_j|$$

- Shrinkage
- **Sparsity**: some fitted coefficients are **exactly** zero

Continuous variable selection

Equivalent Formulation



Equivalent Formulation

$$\begin{aligned} \min_{\beta} \quad & \sum_{i=1}^n \left(y_i - \beta_0 - \sum_{j=1}^p \beta_j x_{ij} \right)^2 \\ \text{subject to} \quad & \sum_{j=1}^p |\beta_j| \leq s \end{aligned}$$

Soft Thresholding

When \mathbf{X} is orthonormal, we can minimize over β componentwise


$$\min_{\beta_j} (\beta_j - \hat{\beta}_j^{\text{ols}})^2 + \lambda |\beta_j|.$$

The solution is

$$\begin{aligned}\hat{\beta}_j^{\text{lasso}} &= \begin{cases} \hat{\beta}_j^{\text{ols}} - \frac{\lambda}{2} & \text{if } \hat{\beta}_j^{\text{ols}} > \frac{\lambda}{2} \\ 0 & \text{if } |\hat{\beta}_j^{\text{ols}}| \leq \frac{\lambda}{2} \\ \hat{\beta}_j^{\text{ols}} + \frac{\lambda}{2} & \text{if } \hat{\beta}_j^{\text{ols}} < -\frac{\lambda}{2} \end{cases} \\ &= \text{sign}(\hat{\beta}_j^{\text{ols}}) \cdot \left(|\hat{\beta}_j^{\text{ols}}| - \frac{\lambda}{2} \right)_+\end{aligned}$$

- Lasso shrinks large coefficients by a constant.
- Lasso truncates small coefficients to zero.

Ridge vs Lasso



Example: Life Expectancy


```
> plot(lmod)
> require(lars)
> data(state)
> statedata = data.frame(state.x77, row.names = state.abb)
> colnames(statedata)
[1] "Population" "Income"      "Illiteracy" "Life.Exp"    "Murder"      "HS.Grad"
[7] "Frost"      "Area"
```

```
> lmod = lars(as.matrix(statedata[,-4]), statedata$Life)
> coef(lmod)
```

	Population	Income	Illiteracy	Murder	HS.Grad	Frost	Area
[1,]	0.00e+00	0.00e+00	0.0000	0.000	0.0000	0.00000	0.00e+00
[2,]	0.00e+00	0.00e+00	0.0000	-0.141	0.0000	0.00000	0.00e+00
[3,]	0.00e+00	0.00e+00	0.0000	-0.203	0.0282	0.00000	0.00e+00
[4,]	1.28e-05	0.00e+00	0.0000	-0.216	0.0308	0.00000	0.00e+00
[5,]	4.90e-05	0.00e+00	0.0000	-0.298	0.0461	-0.00576	0.00e+00
[6,]	4.90e-05	-5.22e-08	0.0000	-0.298	0.0461	-0.00576	0.00e+00
[7,]	4.97e-05	-8.19e-06	0.0000	-0.298	0.0467	-0.00581	-7.79e-09
[8,]	5.18e-05	-2.18e-05	0.0338	-0.301	0.0489	-0.00574	-7.38e-08

Example: Life Expectancy

```
> cvlmod = cv.lars(as.matrix(statedata[, -4]), statedata$Life.Exp)
> which.min( cvlmod$cv )
[1] 66
> cvlmod$index[66]
[1] 0.657
```



Example: Life Expectancy

```
> predict(lmod, s=0.657, type="coef", mode="fraction")$coef
```

Population	Income	Illiteracy	Murder
2.35e-05	0.00e+00	0.00e+00	-2.40e-01

HS.Grad	Frost	Area
3.53e-02	-1.70e-03	0.00e+00

Example: Life Expectancy

```
> g = lm(Life.Exp ~ Population + Murder + HS.Grad +Frost, stated
```

```
> coef(g)
```

(Intercept)	Population	Murder	HS.Grad	Frost
7.10e+01	5.01e-05	-3.00e-01	4.66e-02	-5.94e-03

Example: Life Expectancy

```
# Ridge
> require(MASS)
> g = lm.ridge(Life.Exp ~., statedata, lambda = seq(0, 4, len = 50))
> which.min(g$GCV)
2.7755
> g = lm.ridge(Life.Exp ~., statedata, lambda = 2.7755)
> g
```

	Population	Income	Illiteracy	Murder
7.08e+01	4.13e-05	2.32e-05	-7.89e-02	-2.64e-01
HS.Grad	Frost	Area		
4.60e-02	-5.15e-03	-3.89e-07		

Example: Life Expectancy

```
# AIC  
> g = lm(Life.Exp ~., statedata)  
> step(g, direction = "backward", k = 2)
```

Step: AIC=-28.2

Life.Exp ~ Population + Murder + HS.Grad + Frost

	Df	Sum of Sq	RSS	AIC
<none>			23.3	-28.2
- Population	1	2.1	25.4	-25.9
- Frost	1	3.1	26.4	-23.9
- HS.Grad	1	5.1	28.4	-20.2
- Murder	1	34.8	58.1	15.5

Example: Life Expectancy

Call:

```
lm(formula = Life.Exp ~ Population + Murder + HS.Grad + Frost,  
data = statedata)
```

Coefficients:

(Intercept)	Population	Murder	HS.Grad	Frost
7.10e+01	5.01e-05	-3.00e-01	4.66e-02	-5.94e-03

Lasso

- Useful for high-dimensional data
- Still works when $p \gg n$
- Theoretically guarantees

Simulation Study: High-Dimensional Data

- Response: Y
- Predictors: X_1, X_2, \dots, X_{40}
- 30 samples

```
> dim(data)
```

```
[1] 30 41
```

```
> g = lm(Y ~ ., data)
```

Simulation Study: High-Dimensional Data

```
> coef(g)
(Intercept)  X1          X2          X3          X4          X5
-0.3313      1.8273     -1.3044      8.6732     -2.8432     -1.5281
X6           X7          X8          X9          X10         X11
-3.2323     -3.0793      3.0940     -0.4947     -1.5609     -1.0737
X12          X13         X14         X15         X16         X17
0.0377      3.1824     -3.0936     -0.5792     -0.2540      3.0077
X18          X19         X20         X21         X22         X23
4.3260     -1.1224      2.3879      1.7569      2.8271     -0.5614
X24          X25         X26         X27         X28         X29
2.6789      5.5562     -0.3190     -1.1525     -2.5788      1.4921
X30          X31         X32         X33         X34         X35
NA           NA          NA          NA          NA          NA
X36          X37         X38         X39         X40
NA           NA          NA          NA          NA
```

Simulation Study: High-Dimensional Data

```
> cvlmod = cv.lars(as.matrix(data[,-1]), data$Y)
> which.min( cvlmod$cv )
[1] 24
> cvlmod$index[24]
[1] 0.232
> predict(lmod, s = 0.232, type = "coef", mode = "fraction")$coef
```

X1	X2	X3	X4	X5	X6	X7	X8	X9
0.6951	0.0212	0.0000	0.0000	0.0000	0.0000	0.0000	0.0276	0.0038
X10	X11	X12	X13	X14	X15	X16	X17	X18
0.0000	-0.0953	0.0000	0.0000	0.0523	0.0000	0.0000	0.0000	0.0000
X19	X20	X21	X22	X23	X24	X25	X26	X27
0.0000	0.0000	0.0000	-0.2470	0.0000	0.0000	0.0000	0.0000	-0.2140
X28	X29	X30	X31	X32	X33	X34	X35	X36
-0.3813	0.0000	0.0000	0.0000	0.1900	0.0000	0.0000	0.0000	0.0000
X37	X38	X39	X40					
0.0000	0.0086	0.2222	-0.2335					

Simulation Study: High-Dimensional Data

```
> g = lm.ridge(Y ~., data, lambda = 1)
```

```
> g
```

X1	X2	X3	X4	X5	X6	
1.06e-01	5.55e-01	2.84e-01	4.34e-02	2.17e-02	2.49e-02	2.53e-01
X7	X8	X9	X10	X11	X12	X13
4.12e-02	1.82e-01	2.16e-01	3.13e-02	-2.00e-01	-2.18e-01	2.80e-01
X14	X15	X16	X17	X18	X19	X20
2.90e-01	6.60e-02	-2.50e-01	1.73e-01	-1.66e-01	5.52e-02	2.26e-01
X21	X22	X23	X24	X25	X26	X27
-7.37e-02	-8.42e-02	-9.07e-05	5.40e-02	2.26e-01	1.05e-01	-2.78e-01
X28	X29	X30	X31	X32	X33	X34
-5.33e-01	-1.74e-01	1.48e-02	-2.68e-02	-5.19e-02	3.57e-01	-7.81e-02
X35	X36	X37	X38	X39	X40	
-1.66e-01	-4.16e-02	-9.16e-02	4.18e-01	2.67e-01	-6.45e-01	

Partial Least Square Regression

Partial least squares (PLS) is a method for relating a set of input variables X_1, X_p and outputs $Z_1, \dots, Z_{p'}$. In addition, regress Y over $Z_1, \dots, Z_{p'}$.