

Chapter 10: Shrinkage Methods

Gunwoong Park

Lecture Note

University of Seoul

- Principal components regression
- Partial least squares
- Ridge regression
- Lasso

Principal Components Analysis

Principal Components Analysis (PCA)

PCA:

- Special transformation on predictors
- Useful for high dimensional data
- Solve collinearity issue

Main uses:

- Reduce the dimensionality of the data
- Find linear combinations of predictors that explain the most variation in the data
- Facilitate visualization
- In regression, makes predictors orthogonal to each other


Definition of Principal Components

- Find the u_1 such that $\text{var}(u_1^T X)$ is maximized subject to $u_1^T u_1 = 1$.
- Find the u_2 such that $\text{var}(u_2^T X)$ is maximized subject to $u_2^T u_1 = 0$ and $u_2^T u_2 = 1$.
- Keep finding directions of greatest variation orthogonal to those directions we have already found.

Definitions:

- Vectors u_j are called the PC directions.
- Vectors $z_j = Xu_j$ are called the principal components of X .

Principal Components Analysis (PCA)



Principal Components Analysis (PCA)




Figure 1: PCA: u_1 and u_2

Simulation Study: PCs

```
> prcomp(x)
```

Rotation:


	PC1	PC2
x1	-0.892	-0.451
x2	-0.451	0.892

$$Z_1 = -0.892X_1 - 0.451X_2$$

$$Z_2 = -0.451X_1 + 0.892X_2$$

Simulation Study: PCs

```
> plot(x, pch = 16, xlim = c(-4, 4), ylim = c(-3,3))  
> abline(a = 0, b = 0.451/-0.892, col = "red")  
> abline(a = 0, b = 0.451/0.892, col = "blue")
```



Simulation Study

```
> prX = prcomp(x)
> summary(prX)
```

Importance of components:

	PC1	PC2
Standard deviation	1.210	0.1918
Proportion of Variance	0.976	0.0245
Cumulative Proportion	0.976	1.0000

```
> round(prX$rot[,1],3)
x1      x2
-0.892 -0.451
```

$$Z_1 = -0.892X_1 - 0.451X_2$$

Z_1 explains 97.6% of the both X_1 and X_2

Simulation Study: PCR

```
> lm0 = lm(Y ~ X1 + X2)
> summary(lm0)
```

Coefficients:

Estimate Std. Error t value Pr(>|t|)

(Intercept)	-0.1164	0.1091	-1.067	0.289	
x1	0.8597	0.2020	4.255	4.82e-05	***
x2	1.2688	0.2219	5.717	1.20e-07	***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 1.086 on 97 degrees of freedom

Multiple R-squared: 0.7988, Adjusted R-squared: 0.7946

F-statistic: 192.5 on 2 and 97 DF, p-value: < 2.2e-16

Simulation Study

```
> lm1 = lm(Y ~ Z)
```

```
> summary(lm1)
```

Coefficients:

Estimate Std. Error t value Pr(>|t|)

(Intercept) -0.1092 0.1094 -0.999 0.32

z -1.4875 0.0762 -19.520 <2e-16 ***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 1.09 on 98 degrees of freedom

Multiple R-squared: 0.7954, Adjusted R-squared: 0.7933

F-statistic: 381 on 1 and 98 DF, p-value: < 2.2e-16

Fat Example: Data

- Response: fat
- Predictors: neck, chest, abdom, hip, thigh, knee, ankle, biceps, forearm, wrist

```
> cfat = fat[,9:18]
> prfat = prcomp(cfat)
> dim(prfat$rot)
[1] 10 10
> dim(prfat$x)
[1] 252 10
```

```
> summary(prfat)
```

Importance of components:

	PC1	PC2	PC3	PC4	PC5	PC6	PC7	PC8	PC9
Standard deviation	15.990	4.0658	2.9660	2.0004	1.69408	1.49881	1.30322	1.25478	1.10955
Proportion of Variance	0.867	0.0561	0.0298	0.0136	0.00973	0.00762	0.00576	0.00534	0.00417
Cumulative Proportion	0.867	0.9230	0.9529	0.9664	0.97617	0.98378	0.98954	0.99488	0.99906

Fat Example

```
> round(prfat$rot[,1],2)
```

neck	chest	abdom	hip	thigh	knee
0.12	0.50	0.66	0.42	0.28	0.12

ankle	biceps	forearm	wrist
0.06	0.15	0.07	0.04

```
> prfatc = prcomp(cfat, scale = TRUE)
```

```
> summary(prfatc)
```

Importance of components:

PC1	PC2	PC3	PC4	PC5	PC6	PC7	PC8	PC9	PC10			
Standard deviation			2.650	0.8530	0.8191	0.7011	0.5471	0.5283	0.4520	0.4054	0.27827	0.253
Proportion of Variance			0.702	0.0728	0.0671	0.0492	0.0299	0.0279	0.0204	0.0164	0.00774	0.000
Cumulative Proportion			0.702	0.7749	0.8420	0.8911	0.9211	0.9490	0.9694	0.9859	0.99360	1.000

Fat Example: PCs

```
> round(prfatc$rot[,1],3)
```

neck	chest	abdom	hip	thigh	knee	ankle
0.327	0.339	0.334	0.348	0.333	0.329	0.247

biceps	forearm	wrist
0.322	0.270	0.299

```
> round(prfatc$rot[,2],3)
```

neck	chest	abdom	hip	thigh	knee	ankle
-0.003	-0.273	-0.398	-0.255	-0.191	0.022	0.625

biceps	forearm	wrist
0.022	0.363	0.377

Fat Example: OLS

```
> lmmoda = lm(fat$brozek ~., data = cfat)
> summary(lmmoda)
Coefficients:
Estimate Std. Error t value Pr(>|t|)
(Intercept)  7.228749    6.214309   1.163  0.24588
neck         -0.581947    0.208580  -2.790  0.00569 **
chest        -0.090847    0.085430  -1.063  0.28866
abdom         0.960229    0.071582  13.414 < 2e-16 ***
hip          -0.391355    0.112686  -3.473  0.00061 ***
thigh         0.133708    0.124922   1.070  0.28554
knee         -0.094055    0.212394  -0.443  0.65828
ankle         0.004222    0.203175   0.021  0.98344
biceps        0.111196    0.159118   0.699  0.48533
forearm       0.344536    0.185511   1.857  0.06450 .
wrist        -1.353472    0.471410  -2.871  0.00445 **

Residual standard error: 4.071 on 241 degrees of freedom
Multiple R-squared:  0.7351, Adjusted R-squared:  0.7241
F-statistic: 66.87 on 10 and 241 DF,  p-value: < 2.2e-16
```


Fat Example: PCR

```
> lmodpcr = lm(fat$brozek ~ prfatc$x[,1:2])
> summary(lmodpcr)
Coefficients:
Estimate Std. Error t value Pr(>|t|)
(Intercept)      18.9385      0.3291  57.542 <2e-16 ***
prfatc$x[, 1:2]PC1   1.8420      0.1245  14.800 <2e-16 ***
prfatc$x[, 1:2]PC2  -3.5505      0.3866  -9.184 <2e-16 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 5.225 on 249 degrees of freedom
Multiple R-squared:  0.5492, Adjusted R-squared:  0.5456
F-statistic: 151.7 on 2 and 249 DF, p-value: < 2.2e-16
```

Fat Example

```
> lmodpcr2 = lm(fat$brozek ~ prfatc$x[,1:10])
> summary(lmodpcr2)
```

```
Estimate Std. Error t value Pr(>|t|)
(Intercept)      18.93849    0.25647  73.843 < 2e-16 ***
prfatc$x[, 1:10]PC1    1.84198    0.09698  18.993 < 2e-16 ***
prfatc$x[, 1:10]PC2   -3.55053    0.30126 -11.785 < 2e-16 ***
prfatc$x[, 1:10]PC3    0.25669    0.31374   0.818 0.414067
prfatc$x[, 1:10]PC4    0.54094    0.36652   1.476 0.141273
prfatc$x[, 1:10]PC5    3.72632    0.46973   7.933 8.03e-14 ***
prfatc$x[, 1:10]PC6   -1.48784    0.48642  -3.059 0.002474 **
prfatc$x[, 1:10]PC7    1.94878    0.56859   3.427 0.000716 ***
prfatc$x[, 1:10]PC8   -0.12247    0.63390  -0.193 0.846967
prfatc$x[, 1:10]PC9   -1.71366    0.92351  -1.856 0.064731 .
prfatc$x[, 1:10]PC10 -9.01059    1.01566  -8.872 < 2e-16 ***
```

```
Residual standard error: 4.071 on 241 degrees of freedom
Multiple R-squared:  0.7351, Adjusted R-squared:  0.7241
F-statistic: 66.87 on 10 and 241 DF,  p-value: < 2.2e-16
```

Fat Example: Benefits of PCA

- Orthogonal Predictors
- No collinearity Issue
- Sometimes easy to Interpret

Choice of the number of variables

How to choose the optimal number of variables.

- **Interpretability**: It is important to examine the interpretability of the components and make sure that those providing a interpretable result are retained.
- **Total variance**
- **Cross validation**
- Typically most variation in X can be represented by a few principal components – **Dimension reduction**.
- Can choose k to explain certain **percent of variation**: pick first k so that

$$\sum_{i=1}^k \lambda_i \geq (1 - \alpha) \sum_{i=1}^p \lambda_i$$

- Can look at the **scree plot** and look for a gap in eigenvalues
- More sophisticated methods for estimating **intrinsic dimension** of the data

Choice of the number of variables: RMSE

How to choose the optimal number of variables.

- Training and Test Set
- Root mean square of error (RMSE)


Motivation

- Model Selections and Shrinkage Methods provide good models
- Hard to choose only one optimal model.
- Choose an optimal model based on its performance.

Training and Test Data Set

Method


- Split dataset into two groups .
- **Training Set**: Used to train the model.
- **Test Set**: Used to measure the performance of the trained model.
- **Validation Set**: Used to train tuning parameters. (Extra).



Training and Test Data Set

Why Test Set?

- **Overfitting** Issue: It refers to a model that models the training data too well.



Training and Test Data Set

Why Test Set?

- Overfitting Issue
- Complex models tend to have a good performance in training data.
- A good model in training dataset may not be a good model in new dataset.

How to determine Test Data Set?

- Absolute Random selection: Choose 10 ~ 20% of data set.
- Choose 10 ~ 20% of data set with same proportion of success in both Training and Test data set. (To prevent the test set only contains success or same value of predictors.)


Weakness

- Cannot be performed when sample size is **small**
- Cross-Validation

Cross-Validation

- Often there is insufficient data to create a separate validation set;
 - In this instance, *K-fold cross-validation* is useful.
1. Divide the data into K disjoint subsets.
 2. Use subsets $2, \dots, K$ as *training* data and subset 1 as *validation* data. Compute the PE on subset 1.
 3. Repeat for each subset.
 4. Average the result.

Cross-Validation




Root mean square of error (RMSE)

Definition:

$$\text{RMSE} = \sqrt{\frac{1}{n} \sum_i^n (\hat{y}_i - y_i)^2}$$

Motivation:

- Bias: $y - \mathbb{E}(\hat{y})$
- Variance: $\text{Var}(\hat{y})$



$$y_1 = 10$$

- Small variance case: 95% confidence interval for y_1 : (9.99, 10.01)
→ 95% sure that y_1 is in (9.99, 10.01)
- :Large variance case: 95% confidence interval for y_1 : (-5, 25)
→ 95% sure that y_1 is in (-5, 25)

Root mean square of error (RMSE)

- Mean Square Error = $\text{Bias}(\hat{y})^2 + \text{Variance}(\hat{y})$
- RMSE is a good criterion to choose an optimal model

Simulation Study: Choice of the number of variables

How to choose the optimal number of variables.

- 4 predictors: X_1, X_2, X_3, X_4
- Training set: 40 observations
- Test set: 10 observations

Simulation Study

```
> ran = sample(1:50, replace = F)[1:40]
> train = data[ran,]
> test = data[setdiff(1:50,ran),]
> prx = prcomp(x[ran,])
> summary(prx)
```

Importance of components:

	PC1	PC2	PC3	PC4
Standard deviation	1.3194	0.5817	0.4301	0.21264
Proportion of Variance	0.7538	0.1465	0.0801	0.01958
Cumulative Proportion	0.7538	0.9003	0.9804	1.00000

Simulation Study

```
> lmodpcr2 = lm(y ~ prx$x[,1:2], train)
> z1 = prx$rotation[,1] %*% t( test[,1:4] )
> z2 = prx$rotation[,2] %*% t( test[,1:4] )
> z = rbind(z1, z2)
> ypred = coef(lmodpcr2) %*% rbind(1, z)
> sqrt( mean( (test$y - ypred)^2) )
[1] 1.294111
```

```
> lmodpcr3 = lm(y ~ prx$x[,1:3], train)
> z1 = prx$rotation[,1] %*% t( test[,1:4] )
> z2 = prx$rotation[,2] %*% t( test[,1:4] )
> z3 = prx$rotation[,3] %*% t( test[,1:4] )
> z = rbind(z1, z2, z3)
> ypred = coef(lmodpcr3) %*% rbind(1, z)
> sqrt( mean( (test$y - ypred)^2) )
[1] 1.29453
```

- Interpretation may be easy or difficult
- Sufficiently reduce the number of predictors
- Difficult to decide the number of predictors

Food Analyzer Example

- Response: fat content
- Predictors: 100 channel spectrum of absorbances
- Number of data points: $n = 215$
- Number of predictors: $p = 100$

Prediction Performance

: build a model that predicts well on **future data**.

- Divide the data into two groups: **training data** and **test data**.
- Build the models using the training data and evaluate them on the test data.

Food Analyzer Example Continued

```
> library(faraway)
> data(meatspec)
> dim(meatspec)
[1] 215 101
## Training and test data
> tr = meatspec[1:172,]
> te = meatspec[173:215,]

## Linear model
> g1 = lm(fat ~ ., tr)
> summary(g1)$r.squared
[1] 0.9970196

## Root mean squared error
> rmse = function(x, y) { sqrt(mean( (x - y)^2 ))}
> rmse(g1$fit, tr$fat)
[1] 0.6903167
> ## Prediction
> rmse( predict(g1, newdata=te), te$fat )
[1] 3.814000
```


Food Analyzer Example Continued


```
## AIC
> g2 = step(g1)
> rmse( g2$fit, tr$fat )
[1] 0.7095069
> rmse( predict(g2, newdata=te), te$fat )
[1] 3.590245

## Principal components regression
> library(stats)
> meatpca = prcomp(tr[, -101])
## Square root of the eigenvalues
> round(meatpca$sdev, 3)
[1] 5.055 0.511 0.282 0.168 0.038 0.025 0.014
[8] 0.011 0.005 0.003 0.002 0.002 0.001 0.001
... ...
> matplot(1:100, meatpca$rot[, 1:3], type="l",
xlab="Frequency", ylab="")
```


Food Analyzer Example Continued

```
## Make a scree plot (to choose number of PCs k)
> plot(1:10, meatpca$sdev[1:10], type="l",
xlab="PC number", ylab="SD of PC")
## Fit all PCRs at once and calculate test RMSE for each k
> library(pls)
> pcrg = pcr(fat ~ ., data=tr, ncomp=50)
> rmsmeat = NULL
> for (k in 1:50) {
+   pv = predict(pcrg, newdata=te, ncomp=k)
+   rmsmeat[k] = rmse(pv, te$fat ) }
> plot(rmsmeat, xlab="PC number", ylab="Test RMS")
# scree plot suggestion
> rmsmeat[5]
[1] 3.533628
> which.min(rmsmeat)
[1] 27
> rmsmeat[27]
[1] 1.854858
```


Food Analyzer: First 3 PCs



Food Analyzer: Scree plot



Food Analyzer: Test RMSE




Food Analyzer Example: Cross-validation

```
> pcrg = pcr(fat ~ ., data=tr, ncomp=50,  
> +         validation="CV", segments = 10)  
> rmsCV= RMSEP(pcrg, estimate='CV')  
> which.min(rmsCV$val)  
19  
# Another try  
> pcrg = pcr(fat ~ ., data=tr, ncomp=50,  
> +         validation="CV", segments = 10)  
> rmsCV= RMSEP(pcrg, estimate='CV')  
> which.min(rmsCV$val)  
21  
  
## Plot the RMSE; k=0 is the model with intercept only  
> plot(rmsCV$val, xlab="PC number", ylab="CV RMS")  
## Get test error  
> yfit = predict(pcrg, newdata=te, ncomp=21)  
> rmse(te$fat, yfit)  
[1] 2.214545
```

CV tends to underestimate the real test RMSE but often comes close.

Food Analyzer: Cross-validation



Partial Least Squares

Partial Least Squares

- PCR ignores y when building z 's
- Partial least squares (PLS) chooses z 's that are best at predicting y .
- PLS does not solve a well-defined modelling problem
- Many algorithms for PLS exist
- Also need to select number of components
- No interpretation

Partial Least Squares: Algorithm

Algorithm:

1. Center y , center and standardize each x_j
2. Regress y on each x_j *separately* to get α_j
3. Construct $z_1 = \sum \alpha_j x_j$, which is the first PLS component
4. Regress y on z_1 to get $\hat{\beta}_1$
5. Orthogonalize each x_j with respect to z_1
6. Continue until the final model is fit:

$$\hat{y} = \bar{y} + \hat{\beta}_1 z_1 + \cdots \hat{\beta}_k z_k$$

Remarks:


- Prediction purpose: PLS better
- Explanation purpose: PCR better

Food Analyzer Example

```
> ## Partial least squares
> plsg = plsrf(fat ~ ., data=tr, ncomp=50, validation="CV")
> # plot RMSE estimated by CV
> pls_rmsCV = RMSEP(plsg, estimate='CV')
> plot(pls_rmsCV$val, xlab="PC number", ylab="CV RMS")
> which.min(pls_rmsCV$val)
[1] 14
> ## RMSE on the training data
> dim(plsg$fit)
[1] 172    1   50
> rmse(plsg$fit[, , 14], tr$fat)
[1] 1.952796

> ## RMSE on the test data
> ypred.te = predict(plsg, newdata=te)
> dim(ypred.te)
[1] 43    1   50
> rmse(ypred.te[, , 14], te$fat)
[1] 2.011180
```

Food Analyzer Example (PLS)



Ridge Regression

Ridge Regression

Penalizing the square of the coefficients

$$\min_{\beta} \sum_{i=1}^n \left(y_i - \beta_0 - \sum_{j=1}^p \beta_j x_{ij} \right)^2 + \lambda \sum_{j=1}^p \beta_j^2$$

Assumption:

- Regression Coefficients should not be very large (after standardization).
- A large number of predictors should be considered.
- High collinearity exists.

Ridge Regression

Penalizing the square of the coefficients

$$\min_{\beta} \sum_{i=1}^n \left(y_i - \beta_0 - \sum_{j=1}^p \beta_j x_{ij} \right)^2 + \lambda \sum_{j=1}^p \beta_j^2$$


- The coefficients $\hat{\beta}^{\text{ridge}}$ are shrunk towards zero.
- $\lambda \geq 0$ is a **tuning parameter**.
- λ controls the amount of shrinkage.
- What happens if $\lambda \rightarrow 0$?
- What happens if $\lambda \rightarrow \infty$?

Equivalent Formulation

$$\begin{aligned} \min_{\beta} \quad & \sum_{i=1}^n \left(y_i - \beta_0 - \sum_{j=1}^p \beta_j x_{ij} \right)^2 \\ \text{subject to} \quad & \sum_{j=1}^p \beta_j^2 \leq s \end{aligned}$$

- Explicitly constraint the size of the coefficients.

2-Dim Example



Ridge Regression

When there are many **highly correlated variables**

- $\hat{\beta}^{\text{ols}}$ may have a large coefficient on one variable and a similarly large negative coefficient on its correlated variable (**Unstable**).
- In ridge regression, the size constraint tries to avoid this phenomenon.
- Often **standardize** the predictors first.

Solution

- The solution is

$$\hat{\beta}^{\text{ridge}} = (\mathbf{X}^T \mathbf{X} + \lambda \mathbf{I})^{-1} \mathbf{X}^T \mathbf{y}$$

- $\hat{\beta}^{\text{ridge}}$ is still **linear** in \mathbf{y} .
- $\hat{\beta}^{\text{ridge}}$ is **biased**.

Woodbury matrix identity

Condition

- $\mathbf{A} : p \times p$ matrix, $\mathbf{U} : p \times n$ matrix,
 $\mathbf{C} : n \times n$ matrix, $\mathbf{V} : n \times p$ matrix

Result

- $(\mathbf{A} + \mathbf{UCV})^{-1} = \mathbf{A}^{-1} - \mathbf{A}^{-1}\mathbf{U}(\mathbf{C}^{-1} + \mathbf{VA}^{-1}\mathbf{U})^{-1}\mathbf{VA}^{-1}$
- If $\mathbf{A} = \mathbf{I}_p$, $\mathbf{C} = \mathbf{I}_n$, then

$$(\mathbf{I}_p + \mathbf{UV})^{-1} = \mathbf{I}_p - \mathbf{U}(\mathbf{I}_n + \mathbf{VU})^{-1}\mathbf{V}.$$

Woodbury matrix identity

$$\begin{aligned} & (\mathbf{A} + \mathbf{UCV})[\mathbf{A}^{-1} - \mathbf{A}^{-1}\mathbf{U}(\mathbf{C}^{-1} + \mathbf{VA}^{-1}\mathbf{U})^{-1}\mathbf{VA}^{-1}] \\ &= \{\mathbf{I} - \mathbf{U}(\mathbf{C}^{-1} + \mathbf{VA}^{-1}\mathbf{U})^{-1}\mathbf{VA}^{-1}\} \\ &+ \{\mathbf{UCVA}^{-1} - \mathbf{UCVA}^{-1}\mathbf{U}(\mathbf{C}^{-1} + \mathbf{VA}^{-1}\mathbf{U})^{-1}\mathbf{VA}^{-1}\} \\ &= \{\mathbf{I} + \mathbf{UCVA}^{-1}\} - \mathbf{U}(\mathbf{C}^{-1} + \mathbf{VA}^{-1}\mathbf{U})^{-1}\mathbf{VA}^{-1} \\ &- \mathbf{UCVA}^{-1}\mathbf{U}(\mathbf{C}^{-1} + \mathbf{VA}^{-1}\mathbf{U})^{-1}\mathbf{VA}^{-1} \\ &= \mathbf{I} + \mathbf{UCVA}^{-1} - (\mathbf{U} + \mathbf{UCVA}^{-1}\mathbf{U})(\mathbf{C}^{-1} + \mathbf{VA}^{-1}\mathbf{U})^{-1}\mathbf{VA}^{-1} \\ &= \mathbf{I} + \mathbf{UCVA}^{-1} - \mathbf{UC}(\mathbf{C}^{-1} + \mathbf{VA}^{-1}\mathbf{U})(\mathbf{C}^{-1} + \mathbf{VA}^{-1}\mathbf{U})^{-1}\mathbf{VA}^{-1} \\ &= \mathbf{I} + \mathbf{UCVA}^{-1} - \mathbf{UCVA}^{-1} \\ &= \mathbf{I} \end{aligned}$$

Ridge Regression

Condition

- $\mathbf{X} : n \times p$ matrix

Goal

$$\min_{\beta} \{(y - \mathbf{X}\beta)'(y - \mathbf{X}\beta) + \lambda\beta'\beta\}$$

$$\Leftrightarrow (\mathbf{X}'\mathbf{X} + \lambda\mathbf{I})\beta = \mathbf{X}'y$$

$$\Leftrightarrow \beta = (\mathbf{X}'\mathbf{X} + \lambda\mathbf{I})^{-1}\mathbf{X}'y$$

Ridge Regression

$$\therefore \beta = (\mathbf{X}'\mathbf{X} + \lambda \mathbf{I}_p)^{-1} \mathbf{X}'y$$

$$= (\lambda (\frac{1}{\sqrt{\lambda}} \mathbf{X}' \frac{1}{\sqrt{\lambda}} \mathbf{X} + \mathbf{I}_p))^{-1} \mathbf{X}'y$$

$$= \lambda^{-1} \left\{ \mathbf{I}_p - \frac{1}{\lambda} \mathbf{X}' (\mathbf{I}_n + \frac{1}{\lambda} \mathbf{X}\mathbf{X}')^{-1} \mathbf{X} \right\} \mathbf{X}'y \quad \text{by Woodbury}$$

$$= \lambda^{-1} \{ \mathbf{I}_p - \mathbf{X}' (\lambda \mathbf{I}_n + \mathbf{X}\mathbf{X}')^{-1} \mathbf{X} \} \mathbf{X}'y$$

$\Rightarrow n < p$ 인 경우에 Ridge Regression에서 β 를 추정할 때,

Woodbury를 쓰면 $p \times p$ 대신 $n \times n$ 행렬의 역행렬만 구해도 됨.

Comparison to LSE

$$\hat{\beta}^{\text{ridge}} = (\mathbf{X}^T \mathbf{X} + \lambda \mathbf{I})^{-1} \mathbf{X}^T \mathbf{y}$$

- Even if \mathbf{X} is not full-rank, $(\mathbf{X}^T \mathbf{X} + \lambda \mathbf{I})$ is invertible, thus solve exact collinearity issue.
- $\hat{\beta}^{\text{ridge}}$ has smaller variance than the OLS, thus may have smaller mean square error (MSE).

Shrinkage in Ridge: Special Case

Suppose **orthonormal design** ($\mathbf{X}^\top \mathbf{X} = \mathbf{I}$). Then $\hat{\beta}^{\text{ols}} = \mathbf{X}^\top \mathbf{y}$, and

$$(\mathbf{y} - \mathbf{X}\beta)^\top (\mathbf{y} - \mathbf{X}\beta) = \text{constant} + \sum_{j=1}^p (\beta_j - \hat{\beta}_j^{\text{ols}})^2.$$

Then ridge regression minimizes

$$\sum_{j=1}^p (\beta_j - \hat{\beta}_j^{\text{ols}})^2 + \lambda \sum_{j=1}^p \beta_j^2.$$

Equivalent to the component-wise minimization

$$\min_{\beta_j} (\beta_j - \hat{\beta}_j^{\text{ols}})^2 + \lambda \beta_j^2 \implies \hat{\beta}_j^{\text{ridge}} = \frac{1}{1 + \lambda} \hat{\beta}_j^{\text{ols}}.$$

Shrinkage in Ridge

- Shrink the estimate towards zero by a positive constant less than 1
- $\text{Var}(\hat{\beta}_j^{\text{ridge}}) = \frac{1}{(1+\lambda)^2} \text{Var}(\hat{\beta}_j^{\text{ols}})$.
- $\lambda \uparrow$, shrinkage \uparrow , bias \uparrow , variance \downarrow
- $\lambda \downarrow$, shrinkage \downarrow , bias \downarrow , variance \uparrow .

Simulation Study: Almost Independent Predictors

3 predictors: X_1, X_2, X_3

```
> cor(data[,1:3])
      x1      x2      x3
x1  1.0000  0.0505 -0.142
x2  0.0505  1.0000 -0.143
x3 -0.1425 -0.1428  1.000

> lmod = lm(y ~ x1 + x2 + x3, data)
> lmod$coefficients
(Intercept)      x1      x2      x3
-0.724      1.023      1.299      1.817

> require(MASS)
> #lambda = 0
> lmrid = lm.ridge(y~x1 + x2 + x3, data, lmabda = 0)
> lmrid
(Intercept)      x1      x2      x3
-0.724      1.023      1.299      1.817
```

Simulation Study

```
> #lambda = 0.1
> lmrid2 = lm.ridge(y~x1 + x2 + x3, data, lambda = 0.1)
> lmrid2
(Intercept)      x1          x2          x3
-0.669         1.021         1.297         1.809
>
> #lambda = 10
> lmrid3 = lm.ridge(y~x1 + x2 + x3, data, lambda = 10)
> lmrid3
(Intercept)      x1          x2          x3
3.819         0.848         1.116         1.187
```

Simulation Study: Correlated Predictors

```
> lmod$coefficients
(Intercept)      x1      x2      x3
-0.676      1.132    -0.275    1.779

> cor(data[,1:3])
x1    x2    x3
x1 1.000 0.996 0.999
x2 0.996 1.000 0.998
x3 0.999 0.998 1.000
```

Simulation Study: Correlated Predictors

```
> require(MASS)
> #lambda = 0
> lmrid = lm.ridge(y~x1 + x2 + x3, data, lmabda = 0)
> lmrid
(Intercept)          x1          x2          x3
-0.676          1.132        -0.275         1.779
>
> #lambda = 0.1
> lmrid2 = lm.ridge(y~x1 + x2 + x3, data, lambda = 0.1)
> lmrid2
(Intercept)          x1          x2          x3
-0.720          1.405         0.327         1.195
>
> #lambda = 1
> lmrid3 = lm.ridge(y~x1 + x2 + x3, data, lambda = 1)
> lmrid3
(Intercept)          x1          x2          x3
-0.166          1.260         0.968         0.848
>
> #lambda = 10
> lmrid3 = lm.ridge(y~x1 + x2 + x3, data, lambda = 10)
> lmrid3
(Intercept)          x1          x2          x3
4.85          1.12         1.07         0.74
```

Simulation Study: Comparison to LSE

```
# Generate Training/Test sets
> ran = sample(1:50, replace = F)[1:40]
> train = data[ran,]
> test = data[setdiff(1:50,ran),]

> lmod = lm(y ~ x1 + x2 + x3, train)
> lmrid = lm.ridge(y~x1 + x2 + x3, train, lambda = 0.1)
> lmrid2 = lm.ridge(y~x1 + x2 + x3, train, lambda = 1)

# RMSE
> sqrt(mean((test$y - predict(lmod,test))^2))
[1] 5.4921
>
> ypred = cbind(1, as.matrix(test[, -4])) %*% coef(lmrid)
> sqrt(mean((test$y - ypred)^2))
[1] 5.1544
>
> ypred = cbind(1, as.matrix(test[, -4])) %*% coef(lmrid2)
> sqrt(mean((test$y - ypred)^2))
[1] 5.5591
```

Simulation Study: Normalization

```
> data.scale = scale(data[,1:3])
> data.scale = data.frame(data.scale, y = data$y)
>
> train = data.scale[ran,]
> test = data.scale[setdiff(1:50,ran),]
>
> lmod = lm(y ~ x1 + x2 + x3, train)
> lmrid = lm.ridge(y~x1 + x2 + x3, train, lambda = 0.1)
> lmrid2 = lm.ridge(y~x1 + x2 + x3, train, lambda = 1)

> sqrt(mean((test$y - predict(lmod,test))^2))
[1] 4.8518
>
> ypred = cbind(1, as.matrix(test[, -4])) %*% coef(lmrid)
> sqrt(mean((test$y - ypred)^2))
[1] 4.8484
>
> ypred = cbind(1, as.matrix(test[, -4])) %*% coef(lmrid2)
> sqrt(mean((test$y - ypred)^2))
[1] 4.8587
```

Simulation Study: Another Training/Test Set

```
> ran = sample(1:50, replace = F)[1:40]
> train = data.scale[ran,]
> test = data.scale[setdiff(1:50,ran),]
>
> lmrid = lm.ridge(y~x1 + x2 + x3, train, lambda = seq(0,1, len = 100))
> which.min(lmrid$GCV)
0.080808

> lmrid_GCV = lm.ridge(y~x1 + x2 + x3, train, lambda = 0.0808)
>
> ypred = cbind(1, as.matrix(test[,-4])) %*% coef(lmrid_GCV)
> sqrt(mean((test$y - ypred)^2))
[1] 4.486
>
> lmod = lm(y ~ x1 + x2 + x3, train)
> sqrt(mean((test$y - predict(lmod,test))^2))
[1] 4.7677
```


Simulation Study: Comparison to PCA

```
> X = train[,1:3]
> prx = prcomp(X)
> summary(prx)
Importance of components:

                PC1      PC2      PC3
Standard deviation   1.755 0.06021 0.03225
Proportion of Variance 0.998 0.00117 0.00034
Cumulative Proportion 0.998 0.99966 1.00000

> z = 0.57778 * X[,1] + 0.57848 * X[,2] + 0.57579 * X[,3]
> lmodpcr = lm(train$y ~ z)
> ypred = cbind(1, as.matrix(0.57778 * test[,1] + 0.57848 * test[,2] + 0.57579 * test[,3]))
> sqrt(mean((test$y - ypred)^2))
[1] 4.1143
```

Least absolute shrinkage and selection operator (Chen, Donoho and Saunders 1996; Tibshirani 1996)

$$\min_{\beta} \sum_{i=1}^n \left(y_i - \beta_0 - \sum_{j=1}^p \beta_j x_{ij} \right)^2 + \lambda \sum_{j=1}^p |\beta_j|$$


- Shrinkage
- **Sparsity**: some fitted coefficients are **exactly** zero

Continuous variable selection

Equivalent Formulation

$$\begin{aligned} \min_{\beta} \quad & \sum_{i=1}^n \left(y_i - \beta_0 - \sum_{j=1}^p \beta_j x_{ij} \right)^2 \\ \text{subject to} \quad & \sum_{j=1}^p |\beta_j| \leq s \end{aligned}$$

Ridge vs Lasso



Soft Thresholding

When \mathbf{X} is orthonormal, we can minimize over β componentwise



$$\hat{\beta}_j^{\text{lasso}} = \arg \min_{\beta_j} (\beta_j - \hat{\beta}_j^{\text{ols}})^2 + \lambda |\beta_j|.$$

The solution is

$$\begin{aligned}\hat{\beta}_j^{\text{lasso}} &= \begin{cases} \hat{\beta}_j^{\text{ols}} - \frac{\lambda}{2} & \text{if } \hat{\beta}_j^{\text{ols}} > \frac{\lambda}{2} \\ 0 & \text{if } |\hat{\beta}_j^{\text{ols}}| \leq \frac{\lambda}{2} \\ \hat{\beta}_j^{\text{ols}} + \frac{\lambda}{2} & \text{if } \hat{\beta}_j^{\text{ols}} < -\frac{\lambda}{2} \end{cases} \\ &= \text{sign}(\hat{\beta}_j^{\text{ols}}) \cdot \left(|\hat{\beta}_j^{\text{ols}}| - \frac{\lambda}{2} \right)_+\end{aligned}$$

- Lasso shrinks large coefficients by a constant.
- Lasso truncates small coefficients to zero.

Ridge vs Lasso



Example: Life Expectancy

```
> plot(lmod)
> require(lars)
> data(state)
> statedata = data.frame(state.x77, row.names = state.abb)
> colnames(statedata)
[1] "Population" "Income"      "Illiteracy" "Life.Exp"    "Murder"      "HS.Grad"
[7] "Frost"      "Area"
```




```
> lmod = lars(as.matrix(statedata[,-4]), statedata$Life)
> coef(lmod)
```

	Population	Income	Illiteracy	Murder	HS.Grad	Frost	Area
[1,]	0.00e+00	0.00e+00	0.0000	0.000	0.0000	0.00000	0.00e+00
[2,]	0.00e+00	0.00e+00	0.0000	-0.141	0.0000	0.00000	0.00e+00
[3,]	0.00e+00	0.00e+00	0.0000	-0.203	0.0282	0.00000	0.00e+00
[4,]	1.28e-05	0.00e+00	0.0000	-0.216	0.0308	0.00000	0.00e+00
[5,]	4.90e-05	0.00e+00	0.0000	-0.298	0.0461	-0.00576	0.00e+00
[6,]	4.90e-05	-5.22e-08	0.0000	-0.298	0.0461	-0.00576	0.00e+00
[7,]	4.97e-05	-8.19e-06	0.0000	-0.298	0.0467	-0.00581	-7.79e-09
[8,]	5.18e-05	-2.18e-05	0.0338	-0.301	0.0489	-0.00574	-7.38e-08

Example: Life Expectancy

```
> cvlmod = cv.lars(as.matrix(statedata[,-4]), statedata$Life.Exp)
> which.min( cvlmod$cv )
[1] 66
> cvlmod$index[66]
[1] 0.657
```



Example: Life Expectancy

```
> predict(lmod, s=0.657, type="coef", mode="fraction")$coef
```

Population	Income	Illiteracy	Murder
2.35e-05	0.00e+00	0.00e+00	-2.40e-01

HS.Grad	Frost	Area
3.53e-02	-1.70e-03	0.00e+00

```
> g = lm(Life.Exp ~ Population + Murder + HS.Grad +Frost, statedata)
```

```
> coef(g)
```

(Intercept)	Population	Murder	HS.Grad	Frost
7.10e+01	5.01e-05	-3.00e-01	4.66e-02	-5.94e-03

Example: Life Expectancy

```
# Ridge
> require(MASS)
> g = lm.ridge(Life.Exp ~., statedata, lambda = seq(0, 4, len = 50))
> which.min(g$GCV)
2.7755
> g = lm.ridge(Life.Exp ~., statedata, lambda = 2.7755)
> g
```

	Population	Income	Illiteracy	Murder
7.08e+01	4.13e-05	2.32e-05	-7.89e-02	-2.64e-01
HS.Grad	Frost	Area		
4.60e-02	-5.15e-03	-3.89e-07		

Example: Life Expectancy

```
# AIC
> g = lm(Life.Exp ~., statedata)
> step(g, direction = "backward", k = 2)
```

Step: AIC=-28.2

Life.Exp ~ Population + Murder + HS.Grad + Frost

	Df	Sum of Sq	RSS	AIC
<none>			23.3	-28.2
- Population	1	2.1	25.4	-25.9
- Frost	1	3.1	26.4	-23.9
- HS.Grad	1	5.1	28.4	-20.2
- Murder	1	34.8	58.1	15.5

Example: Life Expectancy

Call:

```
lm(formula = Life.Exp ~ Population + Murder + HS.Grad + Frost,  
data = statedata)
```

Coefficients:

(Intercept)	Population	Murder	HS.Grad	Frost
7.10e+01	5.01e-05	-3.00e-01	4.66e-02	-5.94e-03

- Useful for high-dimensional data
- Still works when $p \gg n$
- Theoretically guarantees

Simulation Study: High-Dimensional Data

- Response: Y
- Predictors: X_1, X_2, \dots, X_{40}
- 30 samples

```
> dim(data)
[1] 30 41
> g = lm(Y ~., data)
```

Simulation Study: High-Dimensional Data

```
> coef(g)
(Intercept) X1      X2      X3      X4      X5
-0.3313     1.8273   -1.3044   8.6732   -2.8432   -1.5281
X6      X7      X8      X9      X10     X11
-3.2323   -3.0793   3.0940   -0.4947   -1.5609   -1.0737
X12     X13     X14     X15     X16     X17
0.0377   3.1824   -3.0936   -0.5792   -0.2540   3.0077
X18     X19     X20     X21     X22     X23
4.3260   -1.1224   2.3879   1.7569   2.8271   -0.5614
X24     X25     X26     X27     X28     X29
2.6789   5.5562   -0.3190   -1.1525   -2.5788   1.4921
X30     X31     X32     X33     X34     X35
NA       NA       NA       NA       NA       NA
X36     X37     X38     X39     X40
NA       NA       NA       NA       NA
```

Simulation Study: High-Dimensional Data

```
> cvlmod = cv.lars(as.matrix(data[,-1]), data$Y)
> which.min( cvlmod$cv )
[1] 24
> cvlmod$index[24]
[1] 0.232
> predict(lmod, s = 0.232, type = "coef", mode = "fraction")$coef
```

X1	X2	X3	X4	X5	X6	X7	X8	X9
0.6951	0.0212	0.0000	0.0000	0.0000	0.0000	0.0000	0.0276	0.0038
X10	X11	X12	X13	X14	X15	X16	X17	X18
0.0000	-0.0953	0.0000	0.0000	0.0523	0.0000	0.0000	0.0000	0.0000
X19	X20	X21	X22	X23	X24	X25	X26	X27
0.0000	0.0000	0.0000	-0.2470	0.0000	0.0000	0.0000	0.0000	-0.2140
X28	X29	X30	X31	X32	X33	X34	X35	X36
-0.3813	0.0000	0.0000	0.0000	0.1900	0.0000	0.0000	0.0000	0.0000
X37	X38	X39	X40					
0.0000	0.0086	0.2222	-0.2335					

Simulation Study: High-Dimensional Data

```
> g = lm.ridge(Y ~., data, lambda = 1)
> g
```

X1	X2	X3	X4	X5	X6	
1.06e-01	5.55e-01	2.84e-01	4.34e-02	2.17e-02	2.49e-02	2.53e-01
X7	X8	X9	X10	X11	X12	X13
4.12e-02	1.82e-01	2.16e-01	3.13e-02	-2.00e-01	-2.18e-01	2.80e-01
X14	X15	X16	X17	X18	X19	X20
2.90e-01	6.60e-02	-2.50e-01	1.73e-01	-1.66e-01	5.52e-02	2.26e-01
X21	X22	X23	X24	X25	X26	X27
-7.37e-02	-8.42e-02	-9.07e-05	5.40e-02	2.26e-01	1.05e-01	-2.78e-01
X28	X29	X30	X31	X32	X33	X34
-5.33e-01	-1.74e-01	1.48e-02	-2.68e-02	-5.19e-02	3.57e-01	-7.81e-02
X35	X36	X37	X38	X39	X40	
-1.66e-01	-4.16e-02	-9.16e-02	4.18e-01	2.67e-01	-6.45e-01	

Summary

- Main reason to use shrinkage: too many predictors or collinearity
- Interpretation is usually lost
- Ridge and lasso give linear models in the original predictors but no inference
- Prediction is usually improved by shrinkage
- All require selecting a **tuning parameter** (number of components for PCR and PLS, λ for ridge and Lasso) – need validation data or cross-validation.